

# ESB with Openshift

COMPANY

GEPARDEC IT SERVICES GMBH

SUPERVISOR

DI (FH) PETER KULCZYCKI

# Content

Infrastructure as Code

Openshift Basics

MicroProfile Specifications


Prototype Design / Implementation

Demo

Evaluations



<http://dev.solace.com>



**OPENSIFT**

<https://www.openshift.com/>

# Infrastructure as Code

## Consistency

All systems are consistent with the definition

## Repeatability

Changes can be repeated with same outcome

## Reproducibility

System and changes can be reproduced

## Disposable

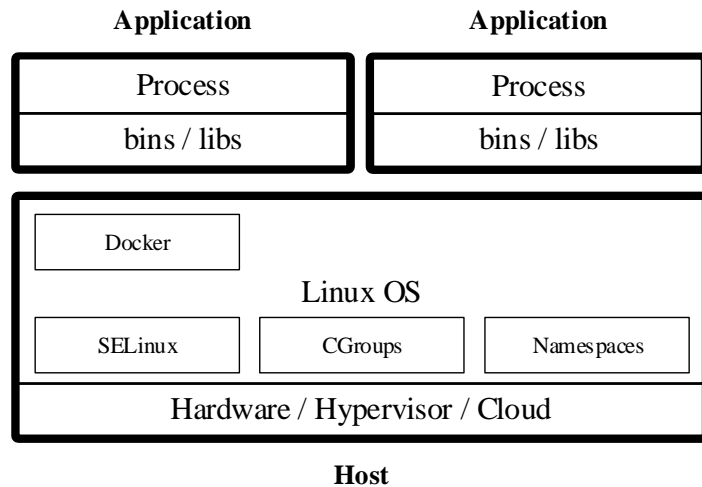
Dispose and Re-Create instead of heavy updates

```
FROM nginx:1.13
MAINTAINER <christoph.ruhsam@gepardec.com>

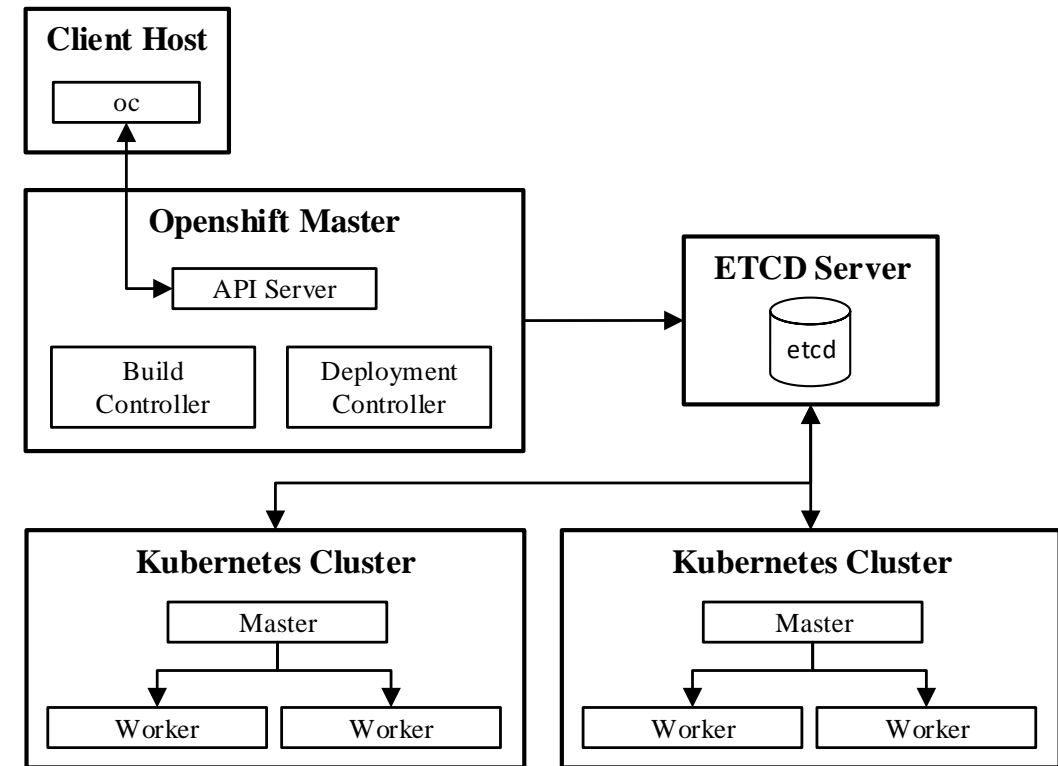
ENV NGINX_RUN_USER swag
ENV NGINX_RUN_GROUP 0
ENV NGINX_HOME_RUN_USER /home/${NGINX_RUN_USER}
ENV DEBIAN_FRONTEND noninteractive

RUN mkdir -p ${NGINX_HOME_RUN_USER} \
    && useradd \
        --home-dir ${NGINX_HOME_RUN_USER} \
        --shell /sbin/nologin \
        -g ${NGINX_RUN_GROUP} \
        -u 1001 \
        ${NGINX_RUN_USER} \
    && apt-get update -y \
    && apt-get install -y curl \
    && apt-get autoremove -y \
    && apt-get autoclean -y \
    && rm -r /var/lib/apt/lists/*
```

# Openshift Basics



**Docker**



**Openshift / Kubernetes Cluster**

# Openshift Basics

## Kubernetes

Namespaces not protected

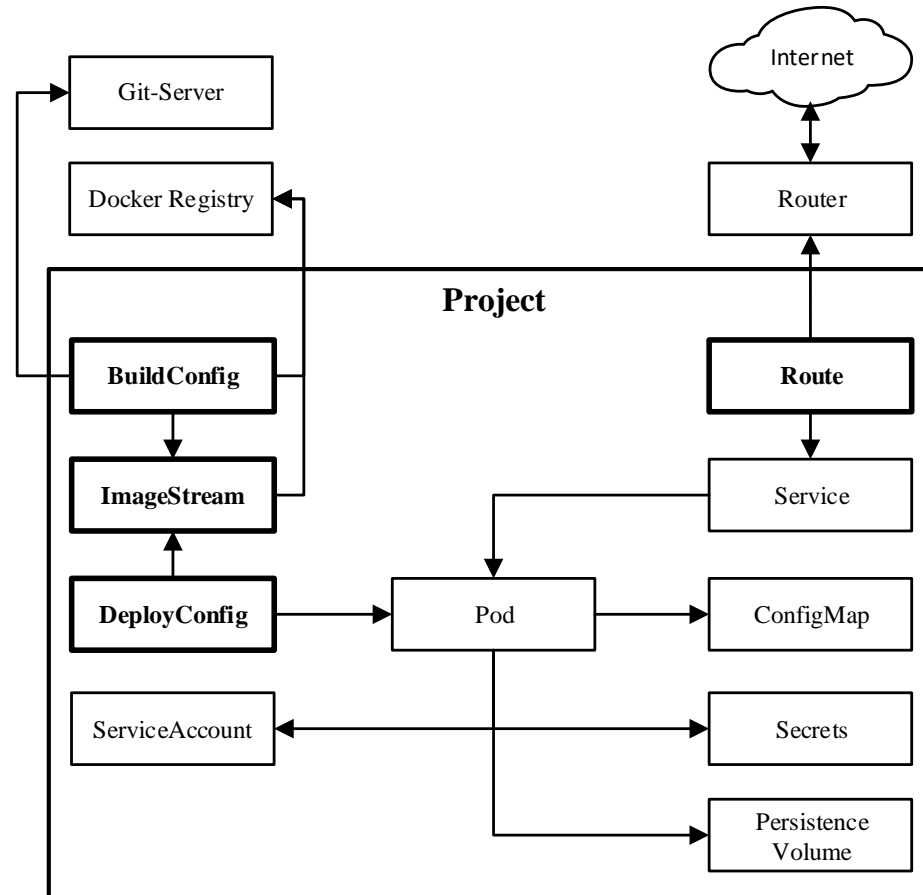
Direct usage of Docker Images

Only cluster state ensured

No rollout behaviors

- Starts Pods
- Stops Pods

No application templates



## Openshift Project

## Openshift

Project isolates Namespaces

Abstraction of Docker Image (IS)

Automated Build -> Deployment

Different Rollout behaviors

- Canary releases
- Blue – Green Releases

Lots of application templates

# Microprofile Specifications

## MicroProfile-Config

ConfigSource, Parameter Injection

## MicroProfile-Fault-Tolerance

Resilience, Retries, Timeouts, Fallbacks

## MicroProfile-Health

Health endpoint ready / liveness probes

## MicroProfile-Metric

Counter, Timer, node state, runtime state

## MicroProfile-OpenTracing

Service / Method call chain tracing




MICROPROFILE  
OPTIMIZING ENTERPRISE JAVA

<https://microprofile.io>



<https://wildfly-swarm.io/>



**HYSTRIX**  
DEFEND YOUR APP

<https://github.com/Netflix/Hystrix>

# Prototype



<https://www.keycloak.org/>




JAEGER

<https://www.jaegertracing.io/>



<https://www.graylog.org/>



MICROPROFILE™  
OPTIMIZING ENTERPRISE JAVA

<https://microprofile.io>



<https://wildfly-swarm.io/>



OPENSIFT

<https://www.openshift.com/>

# Prototype Openshift Project

## Legacy Database

Only accessed by the integration service

## Integration Service

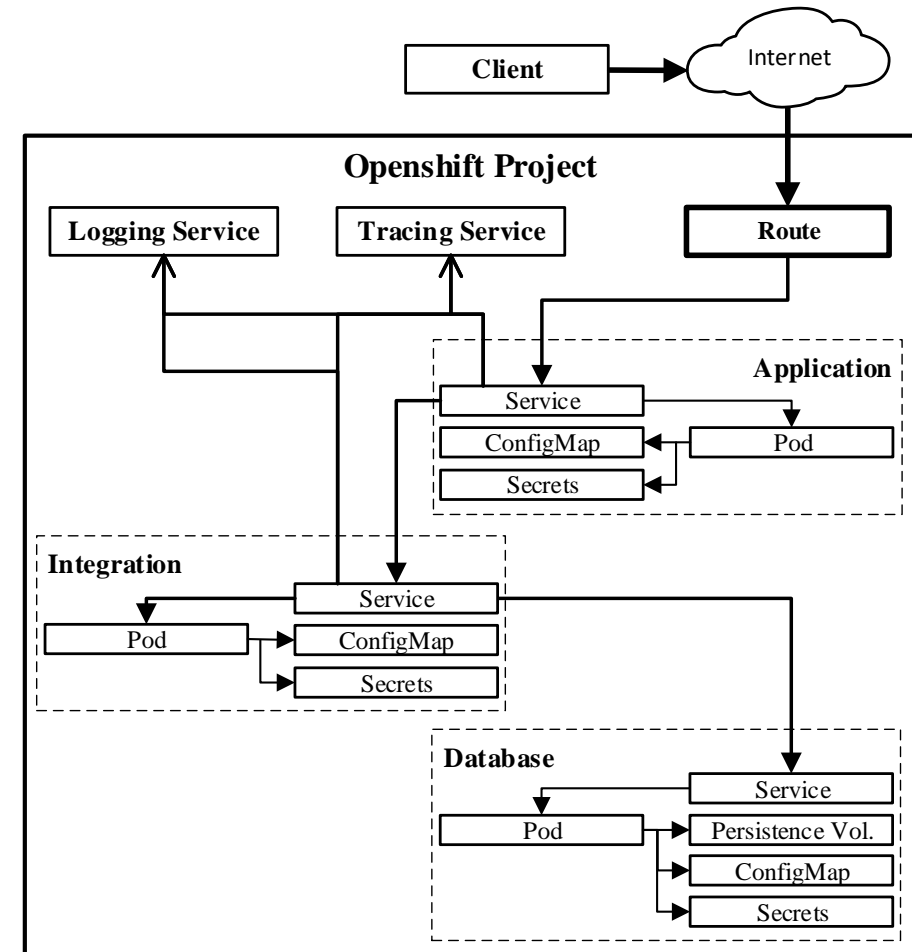
Performs database access and data transformation

## Application Service

Consumes data loaded by the integration service

## Client

Consumes data loaded by the application service





# Prototype Monitoring

## Jaeger (Opentracing)

Tracks spans of calls



## Graylog

Log aggregation over all services



## Keycloak

Oauth2 authentication between services/client-service



## Swagger-UI

Tooling for testing documented REST-API



# Prototype Implementation Security

## Secured Service

```
<dependency>
  <groupId>org.wildfly.swarm</groupId>
  <artifactId>keycloak</artifactId>
</dependency>
```

### 1. Declare Wildfly-Swarm Keycloak fraction

```
volumes:
  - name: keycloak-config
    secret:
      secretName: ${oc.secret-service-integration-keycloak}
```

### 2. Inject secret holding adapter config

```
volumeMounts:
  - name: keycloak-config
    mountPath: ${oc.secret-service-integration-keycloak.dir}
```

### 3. Mount injected secret volume

```
project:
  stage: openshift
  swarm:
    keycloak:
      json:
        path: ${oc.secret-service-integration-keycloak.dir}/keycloak.json
    deployment:
      ${project.build.finalName}.war:
        web:
          login-config:
            auth-method: KEYCLOAK
          security-constraints:
            - url-pattern: /rest-api/customer/*
              roles: [consumer]
            - url-pattern: /rest-api/order/*
              roles: [consumer]
            - url-pattern: /rest-api/admin/*
              roles: [consumer]
```

### 3. Configure Openshift stage security

# Prototype Implementation Security

## Secured Service Consumer

```
volumes:  
  - name: app-config  
    secret:  
      secretName: ${oc.secret-service-app}
```

### 1. Inject secret holding adapter config

```
volumeMounts:  
  - name: app-config  
    mountPath: ${oc.secret-service-app.dir}
```

### 2. Mount injected secret volume

```
swarm:  
  microprofile:  
    config:  
      config-sources:  
        app.secrets:  
          dir: ${oc.secret-service-app.dir}
```

### 3. Define where to load injected configurations

```
@Inject  
@ConfigProperty(name = "keycloak.token-url")  
private String keycloakTokenUrl;  
@Inject  
@ConfigProperty(name = "keycloak.client.id")  
private String keycloakClientId;  
@Inject  
@ConfigProperty(name = "keycloak.client.secret")  
private String keycloakClientSecret;  
  
private ClientCredentialsTokenRequest tokenRequest;  
  
@PostConstruct  
public void postConstruct() {  
    tokenRequest = new ClientCredentialsTokenRequest(new NetHttpTransport(),  
                                                    new JacksonFactory(),  
                                                    new GenericUrl(keycloakTokenUrl));  
    tokenRequest.setClientAuthentication(new ClientParametersAuthentication(keycloakClientId,  
                                                                           keycloakClientSecret));  
}  
  
@Produces  
@OAuthToken  
@Dependent  
@Counted(name = "retrieved-oauth-tokens", monotonic = true)  
@Logging(mdcConfig = Logging.MDCCConfig.GROUP_REST_SECURITY, skipResult = true)  
@Retry(delay = 100L, maxRetries = 5, retryOn = {TokenResponseException.class, IOException.class})  
@Timeout(value = 2L, unit = ChronoUnit.SECONDS)  
String obtainOAuthToken() throws IOException {  
    return tokenRequest.execute().getAccessToken();  
}
```

### 4. Implement Oauth2 token retrieval

# Prototype Implementation Logging

## All Services

```
<dependency>
  <groupId>org.wildfly.swarm</groupId>
  <artifactId>logging</artifactId>
</dependency>

<dependency>
  <groupId>org.jboss.slf4j</groupId>
  <artifactId>slf4j-jboss-logging</artifactId>
  <version>${slf4j.jboss-logging.version}</version>
</dependency>
```

### 1. Declare logging fraction and log implementation

```
@Override
public void filter(ContainerRequestContext requestContext) throws IOException {
    final String tracingId = spanContextInstance.get().toString().split( regex: ":" ) [0]
    scopeInstance.get().span().setTag(MDC_TX_ID, tracingId);
    log.info("Setting MDC transaction id");
    MDC.put(MDC_TX_ID, tracingId);
}
```

### 2. Capture OpenTracing trace id to connect service logs to current trace

```
swarm:
  logging:
    pattern-formatters:
      DEFAULT_LOG_PATTERN:
        pattern: "%d{yyyy-MM-dd HH:mm:ss,SSS} %-5p (%t) [%C{2}] \
                  trace.id=%X{trace.id} | \
                  trace.group=%X{trace.group} - %m%n"
    custom-handlers:
      SYSLOGGER:
        named-formatter: DEFAULT_LOG_PATTERN
        attribute-class: org.jboss.logmanager.handlers.SyslogHandler
        module: org.jboss.logmanager
        properties:
          serverHostname: graylog
          hostname: ${project.build.finalName}
          port: 10514
          protocol: UDP
```

### 3. Configure log formatter and syslog log handler



# Prototype Implementation Tracing

## All Services

```
<dependency>
  <groupId>org.wildfly.swarm</groupId>
  <artifactId>opentracing</artifactId>
</dependency>
```

### 1. Declare OpenTracing fraction

```
env:
- name: "JAEGER_HOST"
  valueFrom:
    secretKeyRef:
      name: ${oc.secret-service-app}
      key: jaeger.host
- name: "JAEGER_PORT"
  valueFrom:
    secretKeyRef:
      name: ${oc.secret-service-app}
      key: jaeger.port
```

### 2. Inject Jaeger host/port from secret into env variables

```
# Necessary because config set via 'swarm.jaeger.*', breaks jaeger integration
JAEGER_SERVICE_NAME: "${project.build.finalName}"
JAEGER_AGENT_HOST: "${env.JAEGER_HOST}"
JAEGER_AGENT_PORT: "${env.JAEGER_PORT}"
JAEGER_REPORTER_LOG_SPANS: "true"
JAEGER_REPORTER_FLUSH_INTERVAL: "1000"
JAEGER_SAMPLER_TYPE: "const"
JAEGER_SAMPLER_PARAM: "1"
```

### 3. Configure OpenTracing fraction

```
@ApplicationScoped
@Traced
@Logging(mdcConfig = Logging.MDCConfig.GROUP_SERVICE)
public class ReportServiceImpl implements ReportService {
```

### 4. Annotate traceable type or method with @Traced

# Prototype DEMO

Home

Overview

Applications

Builds

Resources

Storage

Monitoring

Project

fis-prototype-services

Add to project

thomas.herzo...

> APP SERVICE

<http://app-service-57-fis-prototype-services.cloud.itandtel.at>

Build app-service-s2i, #1 ✔ Complete 5 minutes ago [View Log](#)

> CLIENT SERVICE

<https://management-57-fis-prototype-services.cloud.itandtel.at>  
and 1 other route

Build client-service-s2i, #1 ✔ Complete a minute ago [View Log](#)

> GRAYLOG

<https://graylog-web-57-fis-prototype-services.cloud.itandtel.at>  
and 1 other route

Build graylog, #1 ✔ Complete 5 hours ago [View Log](#)

Build graylog-elastic, #1 ✔ Complete 5 hours ago [View Log](#)

> INTEGRATION SERVICE DB

<http://integration-service-db-57-fis-prototype-services.cloud.itandtel.at>

Build integration-service-db-s2i, #1 ✔ Complete 2 hours ago [View Log](#)

> JAEGER TRACING

Create Route

> KEYCLOAK HTTPS

<http://keycloak-57-fis-prototype-services.cloud.itandtel.at>


> POSTGRESQL PERSISTENT

Create Route


> SWAGGER

<http://swagger-57-fis-prototype-services.cloud.itandtel.at>

Build swagger, #1 ✔ Complete 5 hours ago [View Log](#)



Messages per source



# Evaluations

## Goals of the thesis

### Multistage Configuration

How to configure different kind of services for multiple stages in Openshift, with least of effort ?

### API-Management

How to manage different version of services in Openshift ?

### Transformers and Adapters

How to integrate adapters and transformer services in Openshift ?

### Security

How to secure services in Openshift ?

# Evaluation Multistage Configuration

## MicroProfile-Config

Abstraction from underlying configuration source (ENV, System Props, properties files, custom source)

```
@Inject @ConfigProperty(name="project.developers") private Developer myProp;  
project.developers=[{"name": "Thomas"}, {"name": "Erhard"}]
```

## Wildfly-Swarm project-stages.yml / project-<stage>.yml

Wildfly-swarm configurations support ENV, System Props and Maven Props

```
java -jar -s other-project-stages.yml -Dswarm.project.stage=dev
```

## Openshift ConfigMap / Secret

Application environment provides configuration and secrets and protects them

One project, one stage, one configuration, same binary



# Evaluation API-Mangement

## Need for API Management

Used by multiple clients

External clients could be involved

Are all clients known ?

## Requirement for good API Management

Design a stable API from the beginning

Keep abstraction between internal and external models

Plan your migrations well

Keep only backward compatible as long as needed

Public services need API access tracking



<https://www.benfranklinplumbingmn.com>

# Evaluation API-Mangement

## Path Versioning

/rest-api/customer/v1/list

## Content-Type / Accept Header Versioning

Accept: application/vnd.myapp.user.v1+json;q=0.9

Accept: text/json; version=1;q=0.5

## Query Parameter Versioning

/rest-api/customer/get/1?version=2

## Swagger

Annotations, CodeGen, UI



<https://blog.cloud-elements.com>

# Evaluation Transformers and Adapters

## Global Public Service API *(mostly REST)*

Clients use provided / generated client or integrate it manually. *(Swagger CodeGen)*

Public API must abstract from underlying models, no model pass through

## Data Transformation

Java-Code, Dozer, XSD, JSON, YAML, ... possible in microservice itself

If separately needed, then build up the same way as the other services *(act as proxy service)*

## Adapter

Not needed for internal service hosted on ESB *(Openshift)*

Service in Openshift can act as Adapter for external service

Adapter can be provided for external service

# Evaluation Security

## Openshift Project

Isolates services within a namespace (*Controlled internal/external access*)

## Openshift Secrets

Only reference secret name and property keys, but never values

## Normal JEE-Security

Security Constraints, EJB-Security, Deltaspike-Security, JEE8-Security-API, ...

## Microprofile-Config

No source code has access to configuration source

Thank you for your attention

Any questions ?