

Implementation of an Enterprise Service Bus with OpenShift and Camel

Ing. Thomas Herzog B.Sc



MASTERARBEIT

eingereicht am
Fachhochschul-Masterstudiengang

Software Engineering

in Hagenberg

im Juni 2018

© Copyright 2018 Ing. Thomas Herzog B.Sc

All Rights Reserved

Declaration

I hereby declare and confirm that this thesis is entirely the result of my own original work. Where other sources of information have been used, they have been indicated as such and properly acknowledged. I further declare that this or similar work has not been submitted for credit elsewhere.

Hagenberg, June 1, 2018

Ing. Thomas Herzog B.Sc

Contents

Declaration	iii
Preface	vii
Abstract	viii
1 Introduction	1
1.1 Motivation	1
1.2 Objectives	2
2 Infrastructure as Code	3
2.1 The Need for Infrastructure as Code	3
2.2 Principles of Infrastructure as Code	5
2.2.1 Infrastructures are Reproducible	5
2.2.2 Infrastructures are Disposable	5
2.2.3 Infrastructures are Consistent	6
2.2.4 Actions are Repeatable	6
3 Containerization with Docker	8
3.1 The need for Containerization	8
3.2 Docker	9
3.2.1 Docker Engine	9
3.2.2 Docker Architecture	11
3.2.3 Docker Machine	11
3.3 Virtualization vs. Containerization	12
3.3.1 Virtual Machines	12
3.3.2 Linux Container	13
4 Container as a Service with Kubernetes	15
4.1 The need for Container as a Service	15
4.2 Kubernetes	16
4.2.1 Kubernetes Objects	16
4.2.2 Kubernetes Master	18
4.2.3 Kubernetes Worker	19
5 Platform as a Service with Openshift	20

5.1	The need for Platform as a Service	21
5.2	Openshift	21
5.2.1	Openshift Master	22
5.2.2	Openshift Project	22
6	Enterprise Service Bus	26
6.1	The need for an Enterprise Service Bus	26
6.2	Architecture	26
6.3	JBoss Fuse Middleware	27
6.4	ESB Application	27
6.4.1	Service Component Implementation	27
6.4.2	Service Component Logging	27
6.4.3	Service Component Tracing	27
7	Design ESB in Openshift	28
7.1	Service Architecture	28
7.1.1	Client	29
7.1.2	Service Application	29
7.1.3	Integration Service	29
7.1.4	Database	29
7.2	Service Requirements	29
7.2.1	Technology	30
7.2.2	State Handling	30
7.2.3	Monitoring	30
7.2.4	Configuration	31
7.2.5	Fault Tolerance	31
7.2.6	API Management	31
7.3	Openshift Architecture	32
7.3.1	Project	32
7.3.2	Templates	32
7.3.3	Scripts	32
7.4	Openshift Requirements	32
7.4.1	Monitoring	32
7.4.2	Configuration	32
7.4.3	Secrets	32
8	Implementation ESB in Openshift	33
8.1	Technology	34
8.1.1	JBoss Fuse Integration Services	34
8.1.2	Wildfly Swarm	34
8.1.3	Fabric8	34
8.2	Monitoring	34
8.2.1	Service Implementation	34
8.2.2	Openshift Implementation	34
8.2.3	Jaeger	34
8.2.4	Kibana	34
8.3	Configuration	34

8.3.1	Service Implementation	34
8.3.2	Openshift Implementation	34
8.4	Fault Tolerance	34
8.4.1	Service Implementation	34
8.5	API Management	34
8.5.1	Service Implementation	34
8.6	Openshift	34
8.6.1	Configuration	34
8.6.2	Secrets	34
8.7	Implementation Microservice Architecture	34
8.7.1	Distributed Tracing	35
8.7.2	Distributed Logging	35
8.7.3	API Management	35
8.8	Implementation Openshift Architecture	35
8.8.1	Security	35
9	Evaluation ESB in Openshift	36
9.1	Software Development	36
9.2	Release Management	36
9.3	Security Handling	36
9.4	API Versioning	36
9.5	System Resources	36
	References	37
	Literature	37
	Online sources	37

Preface

Abstract

This should be a 1-page (maximum) summary of your work in English.

Chapter 1

Introduction

1.1 Motivation

Large enterprises work with several independent applications, where each application covers an aspect of a business of the enterprise. In general, these applications are from different vendors, implemented in different programming languages and with their own life cycle management. To provide a business value to the enterprise, these applications are connected via a network and they contribute to a business workflow. The applications have to interexchange data, which is commonly represented in different data formats and versions. This leads to a highly heterogeneous network of applications, which is very hard to maintain.

The major challenge of an IT department is the integration of independent applications into the enterprise application environment. The concept of Enterprise Application Integration (EAI) provides patterns, which help to define a process for the integration of applications into a heterogeneous enterprise application environment. One of these patterns is the Enterprise Service Bus (ESB), which is widely used in the industry [HW08].

Often the term ESB application is used to refer to an ESB, which integrates internal and external hosted applications. But an ESB is a software architectural model, rather than an application. The term could have been established by the usage of middleware such as JBoss Fuse, which provides tooling to integrate applications into an ESB [Red18b]. JBoss Fuse is based on the JBoss Enterprise Application Platform (JBoss EAP), where the applications are integrated in a existing runtime environment.

With the upcoming of cloud solutions such as Platform as a Service (PaaS) it is now possible to move the platform from a dedicated environment to a cloud environment, where each integration service has its own runtime environment rather than joining an existing runtime environment. The concept of Integration Platform as a Service (IPaaS) relies on top of PaaS and enhances a common PaaS solution with the Integration features needed by EAI [DG15; Liu+15a].

Thus, enterprises can reduce the effort in implementing and maintaining an ESB, integrating applications into the ESB and reducing the costs of an ESB by using a consumption based pricing model.

1.2 Objectives

This thesis aims to implement an ESB on Openshift PaaS [Red18c]. Commonly an ESB is implemented with the help of middleware such as JBoss Fuse, which is based on the JBoss EAP. The concepts of PaaS and IPaaS are in general new to the industry, which commonly hosts their integration services in their own data centers, due to the lack of trust for cloud solutions and knowledge about the new approaches such as microservice architecture.

A main focus of this thesis is how applications internal and external can be integrated and managed in the PaaS solution Openshift with the ESB pattern. Before implementing an ESB in a PaaS solution such as Openshift, its necessary to understand the new concepts such as Infrastructure as a Service (IaaS), or containerization with Docker, which are covered in the following chapters. The microservice approach and cloud solutions are becoming more important for the software industry. For instance, Red Hat is currently moving its ESB middleware JBoss Fuse to the cloud, where JBoss Fuse will fully rely on Openshift, and the integration services have to be implemented as microservices. This has huge impact on Red Hats customers, who are used to JBoss Fuse on top of JBoss EAP.

This thesis was commissioned by the company Gepardec IT Services GmbH, a company that is working in the area of Java Enterprise and cloud development. The migration from a monolithic ESB to a microservice structured ESB, which is hosted in a PaaS environment, is a major concern for them. The migration from a monolithic ESB to a microservice structured ESB will be a major challenge for their customers, because microservice architecture and cloud solutions are mostly new to them.

Over the past years, a huge technology dept has been produced by the industry, due to the monolithic architecture and little refactoring work on their applications and hosting infrastructure. It will be hard for them to reduce the produced technology dept, which they will have to, to keep competitive. Gepardec sees a lot of potential for their business and their customers in this new approach of implementing and hosting an ESB.

Chapter 2

Infrastructure as Code

Infrastructure as Code (IaC) is a concept to automate system creation and change management with techniques from software development. Systems are defined in a Domain Specific Language (DSL), which gets interpreted by a tool, which creates an instance of the system or applies changes to it. IaC defines predefined, repeatable routines for managing systems [Kie16]. IaC descriptions are called templates, cookbooks, recipes or playbooks, depending on the tool. In the further course, the IaC definitions will be called templates. The DSL allows to define resources of a system such as network, storage and routing descriptively in a template. The DSL abstracts the developer from system specific settings and provides a way to define the system with as little configuration as possible. The term system is used as a general description. In the context of IaC, a system can be anything which can be described via a DSL.

2.1 The Need for Infrastructure as Code

In the so called iron age, the IT systems were bound the physical hardware and the setup of such a system and its change management were a long term, complex and error prone process. These days, we call such systems legacy systems. In the cloud age, the IT systems are decoupled from the physical hardware and in the case of PaaS they are even decoupled from the operating system [Kie16]. The IT systems are decoupled from the physical hardware and operating system, due to the fact, that cloud providers cannot allow their customer to tamper with the underlying system and hardware. In general, the hardware resources provided by a cloud provider are shared by multiple customers.

With IaC it is possible to work with so called Dynamic Infrastructure Platforms, which provide computing resources, where the developers are completely abstracted from the underlying system. Dynamic infrastructure platforms have the characteristic to be programmable, are available on-demand and provide self service mechanisms, therefore we need IaC to work with such infrastructures [Kie16]. Systems deployed on a dynamic infrastructure platform are flexible, consistent, automated and reproducible.

Enterprises which stuck to legacy systems face the problem that technology nimble competitors can work with their infrastructures more efficiently, and therefore can demand lower prices from their customers. This is due to the IaC principles discussed in Section 2.2 on the following page. Over a short period of time, enterprises will have to move to IaC and away from their legacy systems to stay competitive. The transition process could be challenging for an enterprise, because they lose control over the physical hardware and maybe also over the operating system. Maintaining legacy systems has the effect that someone is close to the system and almost everything is done manually. IaC has the goal to automate almost everything, which requires trust for the cloud providers, who provide the computing resources and the tooling, which provides the automation. A well known problem, which enterprise will face, is the so called Automation Fear Spiral, which is shown in Figure 2.1.




Figure 2.1: Automation Fear Spiral

Because of no trust for the automation, changes are applied manually to the systems and outside the defined automation process. If the system is reproduced, definitions may be missing in the templates, which leads to an inconsistent system. Therefore, enterprises have to break this spiral to fully profit from IaC [Kie16].

When enterprises have moved their legacy systems to IaC, they can not only manage their systems faster, they also can profit from the principles of IaC as discussed in Section 2.2 on the following page. With IaC, systems are less complicated to manage, changes can be applied without fear, and the systems can easily be moved between environments. This provides the enterprises with more space to maneuver, systems can become more complex but still easy to manage, the systems can be defined and

created faster which could lower costs.

2.2 Principles of Infrastructure as Code

The principles of IaC solve the problems of systems of the iron age. In the iron age the creation and maintenance of systems were a long, complicated and error prone process which consumed a lot of resources and time. With the decoupling of the physical hardware from the system, the creation and maintenance of the system has become simple, due to the IaC DSL and tooling.

2.2.1 Infrastructures are Reproducible

With IaC, systems are easy reproducible. It is possible to reproduce the whole infrastructure or parts of it effortlessly. Effortless means, that no tweaks have to be made to the templates or during the reproduction process and there is no need for a long term decision process about what has to be reproduced and how to reproduce it. To be able to reproduce system effortlessly is powerful, because it can be done automatically, consistently and with less risk of failures [Kie16]. The reproducibility of a system is based on reusable templates which provide the possibility to define parameters, which are set for the different environments as shown in Figure 2.2.



Figure 2.2: Schema of a parametrized infrastructure deployment

2.2.2 Infrastructures are Disposable

Another benefit of IaC is that systems are disposable. Disposable means, that systems can be easily destroyed and recreated. Changes made to the templates of a system does not have to be applied on an existing system, but can be applied by destroying and recreating the system. An requirement for a disposable system is, that it is understood that systems will always change. Other systems relying on a disposable system need to address that the system could change at any time. Systems

must not fail because a disposable system disappears and reappears again because of a redeployment [Kie16].

2.2.3 Infrastructures are Consistent

Systems managed with IaC are consistent, because they are defined via a template and all instances are an instance of the template, with the little configuration differences defined by parameters. As long as the system changes are managed by IaC, the system will stay consistent, and the automation process can be trusted.

In Listing 1 an example for an IaC template is shown, which defines a Docker Compose service infrastructure for hosting a Wildfly server instance [Doc18d; Red17]. This system can consistently be reproduced on any environment supporting Docker, Docker Compose and providing values for the defined parameters.

```
1 version: "2.1"
2 services:
3   wildfly:
4     container_name: wildfly
5     image: wildfly:latest
6     ports:
7       - "${EXPOSED_PORT}:8080"
8     environment:
9       - "POSTGRES_DB_URL=${POSTGRES_DB_URL}"
10      - "POSTGRES_DB_NAME=${POSTGRES_DB_NAME}"
11      - "POSTGRES_USER=${POSTGRES_USER}"
12      - "POSTGRES_PASSWORD=${POSTGRES_PASSWORD}"
```

Listing 1: Example for an IaC template for Docker Compose

2.2.4 Actions are Repeatable

Building reproducible systems, means that any action applied to the system should be repeatable. Without repeatability, the automation cannot be trusted and systems wouldn't be reproducible. An instance of a system in another environment should be equal to any other system instance, except for the configurations defined by parameters. If this is not the case, then a system is not reproducible, because it will have become inconsistent [Kie16].

IaC is a concept which makes it very easy to deal with systems in the cloud age. Enterprises can make use of IaC to move their legacy systems to the cloud, where they can profit from the principles of IaC. Nevertheless, before an enterprise can profit from IaC, it has to apply clear structures to their development process, as well as sticking to the principles of consistency and repeatability. For experienced administrators, who are used to maintain systems manually, it could sometimes be hard to understand why they are not supposed to perform any actions on the system manually anymore, nevertheless that a manual change could be performed faster.

Being capable to reproduce a system at any time with no effort, or applying changes on an existing system in a predefined and consistent manner, makes enterprises very flexible and fast. Enterprises will not have to fear future changes in requirements and technologies of their systems anymore.

Chapter 3

Containerization with Docker

Docker is a tool for creating, provisioning and interacting with Linux Containers (LXC) [Doc18f; Lin18a]. LXC are a lightweight version of virtualization, which does not have the resource impact of a full virtualization such as Operating System (OS) virtualization. The differences of LXC and a Virtual Machine (VM) are covered in Section 3.3. Docker has become very popular over the past years, due to the fact, that it made it possible to easily work with LXC. Docker relies strongly on the principles of IaC which has been discussed in Chapter 2. When using Docker, Linux Containers are often referred to as Docker Containers.

Containerization is a key factor when hosting applications in the cloud, because the applications are normally packaged in images and run as containers on the cloud platform. Containerization provides features for a fast, effortless and consistent way of running applications in the cloud, which is discussed in the following Section 3.1.

3.1 The need for Containerization

Containerization is a key factor for cloud platforms such as PaaS, where each application runs in its own isolated environment, called a container. A container is an instance of an image, which represents the initial state of an application. A VM represents a full blown OS, where the OS provides a kernel, which is emulated on the host OS by the Hypervisor. A Hypervisor is a software which can create, run and manage VMs. A container uses the kernel provided by the host OS and therefore there is no need for an emulation. A container does not represent a full blown OS, but still provides features normally provided by an OS such a networking and storage [Sch14].

Containers are faster to create, to deploy and easier to manage compared to VMs. Nevertheless, cloud platforms use virtualization for managing their infrastructure, where the containers run on the provisioned VMs. The usage of containers compared to the usage of VMs can reduce costs for hosting applications. Enterprises can profit from hosting their applications of containers in several ways. Applications hosted in containers need lees resources than applications hosted in VMs, because

there is no virtualized OS and no need for kernel emulation. The creation, deployment and startup of containers are faster, because only the isolated process needs to be started and not a full blown OS. Docker is well supported by Integrated Development Environments (IDEs), which provide support for creating Docker Image definitions (Dockerfiles) and provisioning of Docker Containers on a local or remote environment [Doc18b].

When enterprises have applied IaC to their infrastructure, then the next logical step is to integrate their applications into IaC as well. Applications hosted in containers profit from the IaC principles immutability, reproducibility, repeatability and consistency. Therefore, Docker strongly relies on IaC and provides tooling for automating creation and provisioning of Docker Containers, which is used by PaaS platforms such as Openshift. With Docker, developers define the hosting environment for their applications and not system administrators anymore. Nevertheless, developers can profit from the deep Linux knowledge of system administrators, to define the Docker Images efficiently, to keep them small and secure. The following Section 3.2 will give an overview of the Docker technology, its architecture and artifacts.

3.2 Docker

This section covers Docker, which is the most popular tool to work with LXC. Docker is open source but also provides an enterprise support. The core part of the Docker technology is the Docker Engine, which is discussed in Section 3.2.1. The Docker Engine is the part of the Docker technology that actually runs the containers. The Docker Images are managed in a so called Docker Registry, which is a repository for Docker Images. The most popular Docker Registry is Docker Hub, which is a free service, where anyone can provides Docker Images [Doc18c].

3.2.1 Docker Engine

Figure 3.1 illustrates the Docker Engine architecture hosted on a Linux OS. The Docker Engine is build by layers, where each layer communicates with the layer beneath.



Figure 3.1: Docker Engine architecture

The Docker Engine was initially designed for LXC exclusively but has been ported to Windows. Docker Images and Containers created for Windows OS are not supported on a Linux OS and visa versa. The Docker Images and Containers for a Windows OS differ from those for a Linux OS, but the principles of Docker Images and Docker Containers are the same.

Docker Daemon

The Docker Daemon represents the background process, which creates, runs and manages the Docker Containers on the Docker Host, similar to a VM Hypervisor. The Docker Daemon strongly depends on the kernel of the host OS, therefore incompatibilities could cause the Docker Daemon to fail functioning. The communication with the Docker Daemon is performed via a REST-API, because the Docker Engine is designed as a server client architecture.

REST-API

The REST-API can be exposed via a Unix socket or a network interface, depending on the configuration of the Docker Daemon. If the REST-API is exposed via a network interface, then it is recommended to secure the connection with client certificate authentication. If the Docker Engine and the Docker Client are located on the same host, then commonly the REST-API is exposed via a Unix socket and does not need any special security.

Docker Command Line interface

The Docker Engine provides a Docker Command Line Interface (CLI) for interacting with the Docker Daemon via a Linux shell. The Docker CLI itself communicates with the Docker Daemon via the exposed REST-API. This is the most common way to interact with a Docker Daemon. The Docker CLI provides commands for creating Docker Images and Containers and for provisioning the Docker Containers on the Docker Host.

Docker Images

Docker Images are defined via Dockerfiles, which contain instructions how to build the Docker Image. A Docker Image consists of layers, where each layer represents a state of the file system, produced by a Dockerfile instruction. Each layer is immutable and any change on the file system produces a new layer. Docker Images are hierarchical and can inherit from another Docker Image, which is then called base image. Docker Images support only single inheritance and the base image is defined via the *FROM* instruction as the first instruction in the Dockerfile. Docker Image names have the structure *[namespace]/[name]:[version]* e.g. *library/openjdk:8-alpine*.

Docker Containers

A Docker Container is an instance of a Docker Image, where a new layer is appended, which contains all changes made on the file system by the running process within

the Docker Container. When the Docker Container is deleted, then the appended layer gets deleted as well and all made changes on the file system are lost. A Docker Container keeps running as long as the contained foreground process is running. Without a foreground process the Docker Container stops immediately after it was started. The process running in the Docker Container is isolated from other processes, as well is the file system, the process has access to.

3.2.2 Docker Architecture

The Figure 3.2 illustrates the Docker architecture, which is a client server architecture. The design as a client server architecture is the reason why the communication to the Docker Daemon is performed via the provided REST-API. The Docker Client communicates with the Docker Daemon via the Docker CLI, where the Docker Client can be located on a remote host or on the Docker Host. The Docker Host hosts the Docker Engine, which exposes the REST-API the Docker Client connects to. The Docker Engine managed the Docker Images and Containers located on the Docker Host. The Docker Engine can pull Docker Images from a remote Docker Registry, if a registry has been registered.



Figure 3.2: Docker Architecture

3.2.3 Docker Machine

Docker Machine is a tool for managing local or remote Docker Hosts [Doc18a]. With Docker Machine an administrator can manage multiple Docker Hosts from a main server, without the need to connect to the Docker Host via secure shell (SSH). The Docker Machine CLI provides all commands necessary for managing Docker Hosts. Docker Engine provisions Docker Containers on a Docker Host and Docker Machine provisions Docker Hosts, in particular Docker Engines installed on docker Hosts. With Docker Machines a network of Docker Hosts can be managed, which is used by cloud platforms such as Openshift to manage Docker Engines on the nodes within the Openshift cluster.

3.3 Virtualization vs. Containerization

Before LXC the industry made heavy use of operating system (OS) virtualization to isolate their environments and applications. A VM is managed by a Hypervisor, which is software, which can create, run and manage VMs. The VM provides resources such as network and storage for the application, which is managed by the virtualized OS. Nevertheless, an VM represents a full blown OS, which itself has a resource need which adds to the resource needs of the hosted application. LXC on the other hand are a kernel technology, which provides resources such as network and storage to the application as well, but without the need of virtualized OS.

3.3.1 Virtual Machines

A Virtual Machine is an instance of a Virtual Machine Image (VMI), which is managed by a Virtual Machine Monitor (VMM), which is also referred to as the Hypervisor. The actual difference between a VMM and a Hypervisor is where the software is installed on. If the software is directly installed on the Hardware, then the software is called a Hypervisor, if its installed on the Host OS then its called a VMM. The VM abstracts an Guest OS from the Host OS, in particular from the underlying hardware. A VM contained Guest OS is not bound to the underlying hardware, because the Hypervisor performs a kernel emulation, which allows to virtualize any Guest OS on any hardware, if the hypervisor supports it. The following Figure 3.3 illustrates the architecture of a virtualization system.



Figure 3.3: Architecture of virtualized applications

Glauber Costa's started the abstract of his talk at the LinuxCon 2012 with the humorous note *"I once heard that Hypervisors are the living proof of operating system's incompetence"*. With this note he expressed that OS weren't able to provide proper isolation for applications and therefore the industry started to provide an OS instance for each application [Cos12]. This has been overcome with the upcoming of LXC, which provide the proper isolation of applications on the same OS, which made the need for an OS instance for each application obsolete.

3.3.2 Linux Container

The upcoming of LXC has eliminated the shortcoming to not be able to isolate applications properly of the Linux OS, which lead to using OS virtualization to isolate applications. LXC provide the feature of isolating applications running on the same OS, without the need of a kernel and hardware emulation as it is done with OS virtualization. As illustrated in Figure 3.4, the application process, binaries and libraries are bundled into the container and are isolated from other containers. Each container gets a portion o the global resources such as CPU cycles and memory assigned and cannot consume more as it has been assigned to. Without LXC it is possible that one process takes over the system resources and other processes get into state of starvation, which lead to need of OS virtualization.



Figure 3.4: Architecture of containerized applications

The two most important kernel features underlying LXC are *Cgroups* and *Namespaces*. These two kernel features provide the resource control and isolation needed for application isolation and prevention of process starvation.

Cgroups

Cgroups stands for control groups and Cgroups provide the ability to aggregate processes, their child processes and threads within theses processes to groups managed in a tree structure. Each group gets a portion of the global resources such as CPU time, memory, I/O and network assigned, where its guaranteed that a group and its managed processes cannot consume more resources as the group has been assigned to. Each application hosted in a container is assigned to a group, where an application cannot steal resources from another application anymore, because the resource assignments of an group managed by Cgroups prevents this from happening [Cor14; Heo15; Men18].

Namespaces

Cgroups manage how many resources can be used by processes in a group and

namespaces manage the view of the system to processes. A container is managed in a namespace and therefore it has a limited view of the system such as networks and Process IDs (PIDs), depending on the configuration of the namespace the container is part of. Namespaces are a fundamental concept of LXC, and namespaces provide the isolation of a container [Cor14; Lin18b].

Docker has made the usage of LXC simple, but it is very hard to maintain a large set of Docker Containers (>100) via the Docker CLI, or to implement and maintain a cluster of Docker Hosts with Docker Machine. To much would have to be scripted manually, which would fast become very hard to maintain. Additionally, Docker does not provide any workflow for deployment and scaling of Docker Containers, and also does not ensure that a desired state of the containers is met. For a local development or a small set of containers the Docker CLI, Docker Compose and Docker Machine are suitable, but when it comes to large dynamic infrastructures with a large set of Docker Containers to maintain, then container orchestration platforms like Kubernetes, which is discussed in Chapter 4 on the next page, will have to be used [Doc18e; Kub18d].

Chapter 4

Container as a Service with Kubernetes

Container as a Service (CaaS) is a term introduced by cloud providers, which provide a cloud based on demand container environment. But CaaS is more than just an on demand container environment like Docker, it provides orchestration and monitoring tooling for containers, and additionally CaaS is considered to be a model for IT organizations and developers how they can ship and run their applications anywhere. There are multiple CaaS providers on the market, but the most popular CaaS providers are Azure Container Service, Amazon Elastic Container Service for Kubernetes (Amazon EKS) and Google Kubernetes Engine, where they bring in their own flavor of CaaS but all of them use Kubernetes beneath [Ama18a; Goo18b; Kub18d; Mic18b].

Kubernetes is a container orchestration platform for automating deployments, scaling and operation of containers across a Kubernetes Cluster of Kubernetes Worker-Nodes. Kubernetes has been invented by Google and is open source since 2015 and managed by the Cloud Native Computing Foundation, where the Cloud Native Computing Foundation is under the umbrella of the Linux Foundation. Kubernetes has become the most popular container orchestration platform on the market and is used by many CaaS and PaaS providers [Clo18a].

4.1 The need for Container as a Service

Enterprises and developers are facing the need to dynamically apply to workloads and to roll out new version of their services fast. For applying dynamically to workloads a dynamic infrastructure is necessary to scale services up if the workload increases and to scale services down when the workload decreases, which is non trivial to be handled manually. Rolling out new versions requires a well defined workflow which specifies the roll out behavior, which also is non trivial to handle. For such uses cases a container orchestration platform like Kubernetes can be used, which provides workflows for roll out and support for scaling containers along with many other features. Kubernetes makes it possible to effortlessly manage complex service infrastructures, service scaling and the roll out of services. Thus, complex service infrastructures become simple to implement and manage.

Kubernetes uses IaC, which has been discussed in Chapter 2 on page 3, and therefore provides all of the principles of IaC as discussed in Section 2.2 on page 5. Kubernetes provides a DSL, which allows to specify the desired state of the Kubernetes Cluster such as running containers, container replicas and provided container resources such as RAM, CPU and network. Kubernetes automatically ensures that the state of the Kubernetes Cluster meets its specification. Thus, the developers have only the need to specify the desired state of their Kubernetes Cluster. Kubernetes provides enterprises an infrastructure for their services, which is effortlessly to specify and maintain, because of the automation tooling provided by Kubernetes. This makes it easy to modify the infrastructure at any time, which allows enterprises to apply fast to new requirements.

4.2 Kubernetes

Kubernetes is a platform to orchestrate containers in a cluster, where the Kubernetes Cluster-Nodes can be placed in the cloud or on a dedicated servers. Kubernetes is designed as a client server architecture and a master slave architecture. One node in the Kubernetes Cluster acts as the Kubernetes Master, which is discussed in Section 4.2.2 on page 18, and the other nodes in the Kubernetes Cluster act as the Kubernetes Workers, which are discussed in Section 4.2.3 on page 19. The Figure 4.1 illustrates the architecture of a Kubernetes Cluster.




Figure 4.1: Architecture of a Kubernetes Cluster

4.2.1 Kubernetes Objects

Kubernetes Objects are persistent objects in the Kubernetes System, and the Kubernetes Objects describe the state of the Kubernetes Cluster. The Kubernetes Cluster

ensures that the state of the cluster meets the state specified by the Kubernetes Objects. The developers don't have to manually perform actions in the Kubernetes Cluster, they just have to modify the specification of the state of the Kubernetes Cluster and the Kubernetes Clusters itself will ensure that the new state is applied on the Kubernetes Cluster. The following sections will describe some of the common used Kubernetes Objects. The overview of all Kubernetes Objects is covered by the Kubernetes API reference documentation [Kub18a].

Pod

A Pod is a group of one or more containers which are managed together. A Pod specification contains the specification for each container in the group. All of the containers of a Pod are always scheduled on the same Kubernetes Worker and will be deployed, started and stopped as a single unit. In a pre-container world all of the applications represented by the containers would have been hosted on the same physical machine. A Pod allows to bundle containers together which are acting as a single service, for instance a web application container with a caching container [Kub18c].

Service

A service is an abstraction which defines a set of Pods and policies how to access them. The connection to the Pod via the service abstraction is handled by the Kubernetes Proxy. The service abstraction is necessary because a Pod can be hosted on any Kubernetes Worker within the Kubernetes Cluster, and the Pod will therefore get a random IP assigned which makes it impossible to address the Pod directly. If multiple replicas of a Pod are running, then the service will connect to a Pod of the replica set, depending on the chosen algorithm [Kub18e].

Secret

A secret is an abstraction to manage sensitive data, which is consumed by containers. A secret holds sensitive data and hides it behind a name. The secret can be referenced by a container specification by its name. A referenced secret will be injected into the container either as an environment variable or a file. Developers reference secrets in the specifications by their name, and only the referencing containers can access the sensitive data the secret holds.

ConfigMap

A configuration map is similar to a secret, but is intended to hold non sensitive data. A configuration map is meant to hold configurations such as logging configuration, which is consumed by containers. Configuration map also hide the data behind a name which can be referenced by a specification, and can they can be injected into containers the same way as secrets are. Configurations can be replaced during the container is up and and will be re-injected when the container restarts.

4.2.2 Kubernetes Master

The Kubernetes Master is the master node in the Kubernetes Cluster. It is responsible for managing the Kubernetes Worker-Nodes and containers running on those nodes. The Kubernetes Master exposes a REST-API via the clients can interact with the cluster. The node hosting the Kubernetes Master should be exclusively for the Kubernetes Master. The following sections briefly introduce the Kubernetes Master-Components, which are responsible for managing the Kubernetes Cluster [Kub18b].

Kubernetes CLI (kubectl)

Kubectl is the CLI of Kubernetes, which provides an interface to manage the Kubernetes Cluster and to manage Pods running on the Kubernetes Worker-Nodes. Kubectl is similar to the Docker CLI, but does not support direct interacting with Docker. Kubectl interacts with the Kubernetes Cluster via a REST-API exposed by the Kubernetes Master API-Server. Kubectl can be used from any client machine which can connect to the cluster without any special setup.

Distributed Key-Value Store (etcd)

Etcd is a distributed key-value store and provides a reliable way for sharing data within a cluster. It is the key component for the communication between the Kubernetes Master and the Kubernetes Worker-Nodes. The Kubernetes Master provides configuration for the Kubernetes Nodes and retrieves state information from the Kubernetes Worker-Nodes [Cor18].

Kubernetes API-Server (kube-apiserver)

The Kubernetes API-Server exposes the interface for interacting with the Kubernetes Cluster and is located on the Kubernetes Master. It represents the frontend of the Kubernetes Cluster and provides all necessary API to manage the cluster and the Pods running on it.

Kubernetes Scheduler (kube-scheduler)

The Kubernetes Scheduler watches the Kubernetes Cluster for newly created Pods and assigns the Pods to a Kubernetes Worker-Node. The Kubernetes Scheduler decides which Kubernetes Worker-Node is suitable for the Pod. Multiple factors are taken into account for scheduling decisions such as individual specifications, resource requirements, available resources and hardware/policy/software constraints.

Kubernetes Controller Manager (kube-controller-manager)

The Kubernetes Controller Manager is responsible for the managing of the different controllers. A Kubernetes Controller is running in a loop and ensures that the state of the system is valid, depending on the controller type. For instance, the replication controller ensures the correct number of Pods for each replication controller object

within the Kubernetes Cluster. Kubernetes provides a set of controllers such as a replication controller, node controller, endpoint controller and service account controller.

4.2.3 Kubernetes Worker

The Kubernetes Worker is a node within the Kubernetes Cluster which acts as the slave node which hosts the Pods and is managed by the Kubernetes Master. The Kubernetes Worker can be a VM or a physical machine depending on the Kubernetes Cluster setup. It contains the Kubernetes Runtime-Environment and Docker. The following sections briefly introduce the Kubernetes Worker-Components, which are responsible for running the Pods on the Kubernetes Worker-Node [Kub18b].

Kubernetes Agent (kubelet)

The Kubernetes Agent is a process running on the Kubernetes Worker-Nodes which interacts with the Kubernetes Master via the Kubernetes API-Server. The Kubernetes Agent ensures that the containers are running in a Pod as specified by the provided Pod specifications. The Pod specifications can be provided by a file in a specific directory (gets periodically checked), or via the Kubernetes API-Server.

Kubernetes Network-Proxy (kube-proxy)

The Kubernetes Network-Proxy manages the networks defined by the specifications and reflects the services which are bound to a Pod. It can perform simple TCP and UDP forwarding and can be connected to multiple backends. Any communication of a Pod to another Pod or to the Internet is handled by the Kubernetes Network-Proxy.

Container Runtime

The container runtime is the software responsible for running the containers on the Kubernetes Worker. Kubernetes supports multiple container runtimes, but usually its Docker which has been discussed in Chapter 3 on page 8.

Kubernetes provides all features to implement a dynamic scalable service infrastructure such as workflows for rolling out services, replica management, secret and configuration management, which enterprises can profit from. Secrets are protected from being accessed by the developer and configurations can be applied without building a new service artifact. Kubernetes enhances Docker with orchestration tooling necessary to run large scale dynamic containerized service infrastructures. Nevertheless, sometimes even Kubernetes is not suitable enough for some use cases, which can be overcome with PaaS platforms like Openshift, which is discussed in the following Chapter 5 on the next page.

Chapter 5

Platform as a Service with Openshift

Platform as a Service (PaaS) is a cloud service which provides an on demand platform for building, deploying and running containerized applications in the cloud. PaaS can be seen as an enhancement of CaaS, which has been discussed in the former Chapter 4 on page 15. A PaaS platform does not only provide a container runtime for running containers in the cloud but also tooling for building, deploying and monitoring of containerized applications as well as security mechanisms for securing those applications. There are multiple PaaS providers on the market but the most popular PaaS providers are RedHat Openshift Online, Microsoft Azure Cloud Services, Google App Engine and AWS Elastic Beanstalk. They all bring in their own flavor of PaaS but they all provide similar features necessary by an PaaS platform [Ama18b; Goo18a; Mic18a; Ope18a].

PaaS providers usually provide templates for the major programming languages and application servers, and integration to other cloud services as well. External cloud services of the same vendor are usually better supported than cloud services of other vendors. This is normal, because cloud providers want the developers to use their service over the services of the competition. What all PaaS providers have in common is the consumption based pricing model, where only the consumed physical resources have to be paid for.

IPaaS can be seen as an enhancement of PaaS which is suitable for implementing an ESB which is discussed in Chapter 6 on page 26. IPaaS enhances an ordinary PaaS platform by providing tooling for integrating external service effortlessly, via a low/no code platform, where services can be integrated via an UI, rather than by implementing source code. RedHat JBoss Fuse 7 is an example for an IPaaS platform which will replace JBoss Fuse 6.x in the near future [Liu+15a; Liu+15b; Red18a].

Openshift Origin is an open source PaaS platform, which has been released in April 2012 and is the upstream project for Openshift. Before Openshift 3 (Jun 2013), Openshift used its own container runtime and orchestration tooling, which since Openshift 3 have been replaced by Docker and Kubernetes, because of its popularity and general availability. Openshift is the only major PaaS platform of the formerly

noted ones which can be self hosted or hosted by a local provider. The other formerly noted PaaS providers such as Microsoft Azure are only available as a cloud service hosted in the vendors data centers [Ope18b].

5.1 The need for Platform as a Service

As mentioned in Section 4.2.3 on page 19, there are some use cases where Kubernetes or in particular CaaS is not suitable anymore. CaaS is suitable if its used by developers, but not for persons without any deep knowledge of Docker and Kubernetes. This is where PaaS platforms come into place, which provide a web console and a template mechanism, which can be used by non-developers. Developers specify templates for the provided services which contains all technical parts of a service infrastructure and non-developers provide values for the exposed parameters which are non-technical, and the PaaS platform instantiates the template and deploys the service infrastructure automatically.

Enterprises can profit from PaaS platforms by defining templates for services they provide for their departments, partners or customers, who can create an instance of a provided service on demand, and destroy it if not needed anymore. PaaS platforms provide a self service console, where services can be created, managed and destroyed effortlessly without the need to understand the underlying technology. The self service console could be implemented by enterprises for their specific use cases, where the self service console interacts with PaaS platform via its exposed API.

PaaS platforms like Openshift usually provide an integration in a Continuous Integration / Continuous Deployment (CI/CD) workflow, which allows to automatically build and deploy new service releases in the PaaS platform automatically via web hooks. Therefore, the PaaS platforms are integrated in the whole software life cycle. This decreases the effort of the developers to interact with the cloud platform and provide additional automation.

5.2 Openshift

Openshift is a open source PaaS platform, which uses Docker and Kubernetes for the Docker Container orchestration. Openshift is designed as a client server architecture and a master slave architecture, the same way as a Kubernetes Cluster, which has been discussed in Section 4.2 on page 16. An Openshift Cluster can contain multiple Kubernetes Clusters which are managed by a Openshift Master-Node, which is discussed in Section 5.2.1 on the following page, which manages the Kubernetes Master-Nodes. Openshift provides Openshift Projects, which are discussed in Section 5.2.2 on the next page, which place all defined resources in a Kubernetes namespace, and which are isolated form each other. The following Figure 5.1 on the following page illustrates the architecture of an Openshift Cluster [Ope14; Ope18c].



Figure 5.1: Architecture of a OpenShift Cluster

5.2.1 OpenShift Master

The OpenShift Master-Node manages the Kubernetes Master-Nodes of the Kubernetes Clusters the OpenShift Cluster contains. The OpenShift Master exposes a REST-API via the clients can interact with the OpenShift Cluster. Therefore that OpenShift is placed on top of Kubernetes, the OpenShift Master-Node acts similar as a Kubernetes-Master-Node, which has been discussed in Section 4.2.2 on page 18. Additionally OpenShift provides features Kubernetes does not, such as a role and group based security model for isolating the Kubernetes Namespaces via OpenShift Projects and controllers for managing the additional OpenShift Objects. The following Section 5.2.2 discusses OpenShift Projects, which are the main feature provided by OpenShift.

5.2.2 OpenShift Project

An OpenShift Project represents a Kubernetes namespace, where all resources of an OpenShift Project are located. An OpenShift Project provides the isolation and security Kubernetes Namespaces do not provide. The Figure 5.2 on the next page illustrates the OpenShift Project-Architecture, its contained Objects and their dependency to each other. The bold marked objects within the Project representing

the Openshift Objects which are provided by Openshift.



Figure 5.2: Architecture of a Openshift Project

Openshift Objects are persistent objects in the Openshift System, and the Openshift Objects describe the state of the Openshift Cluster. This behavior has been inherited from the underlying Kubernetes System as discussed in Section 4.2.1 on page 16. The following sections briefly introduce the new Objects provided by Openshift.

BuildConfig

A Build Configuration specifies the way how a Docker Image is built on the Openshift platform. The built Docker Image is pushed into the Openshift internal Docker Registry. Openshift Build Configurations support the following listed strategies:

- The *Source-to-Image (S2I)* strategy is the build strategy which builds a Docker Image from source code.
- The *Docker* strategy is the build strategy which builds a Docker Image from a Dockerfile.

- The *Custom* strategy is the build strategy which build a Docker Image with a custom implemented build mechanism.
- The *Pipeline* strategy is the build strategy which performs a Jenkins pipeline build on a Jenkins build server.

The necessary resources for the particular build strategy are provided via a git repository, and a Build Configuration can be triggered by an external service such as Github via a web hook [Ope18d].

ImageStream

An Image Stream and its Image Stream-Tags are an abstraction of the actual used Docker Image and an Image Stream uses the same naming convention as Docker Tags (E.g *myproject/app:1.0*), where

- *myproject* represents the Image Stream namespace,
- *app* represents the Image Stream name and
- *1.0* represents the Image Stream-Tag.

An Image Stream-Tag references the actual Docker Image by its tag. Once the Docker Image has been imported, it will not be automatically pulled again unless the Image Stream-Tag has the name *latest* which causes Openshift to always to pull the referenced Docker Image.

A Docker Image can be updated in a Docker Registry, which would break the consistency principle, because it wouldn't be the same Docker Image as used before the update. An Image Stream or in particular the Image stream-Tag prevents this, by referencing the actual Docker Image instance instead of only referencing the Docker Image by its tag. This approach makes the Docker Image immutable within a Open-shift Project, unless the latest version is explicitly defined.

DeployConfig

A Deployment Configuration specifies how a deployment of an Pod has to be performed. A Deployment Configuration allows to specify the Kubernetes life cycle hooks pre-hook or post-hook, which are used to configure the deployed Pod before its process has started (pre-hook) or after its process has started and is ready (post-hook). Deployment Configurations support the following listed deployment strategies:

- The *Rolling* strategy is the deployment strategy which waits for the new deployment to be ready before the old deployment gets removed.
- The *Recreate* strategy is the deployment strategy which removes the old deployment when the new deployment gets started.
- The *Custom* strategy is the deployment strategy which performs the deployment by a custom implementation.

Route

A Route exposes a Service with a host name to an external network (mostly the Internet), so that it can be reached by its host name from clients located outside of the Openshift Cluster. The Route is deployed on a Openshift Router, which performs the routing between the external network and the connected Service. A Route can be secured with TLS, where the certificates of the Openshift Cluster can be used or the certificate can directly be defined in the Route definition.

Chapter 6

Enterprise Service Bus

An Enterprise Service Bus (ESB) is a architectural pattern which describes a distributed computing architecture, where distributed services are connected to each other via the ESB. The ESB pattern is part of the Service Oriented Patterns (SOA)

6.1 The need for an Enterprise Service Bus

6.2 Architecture

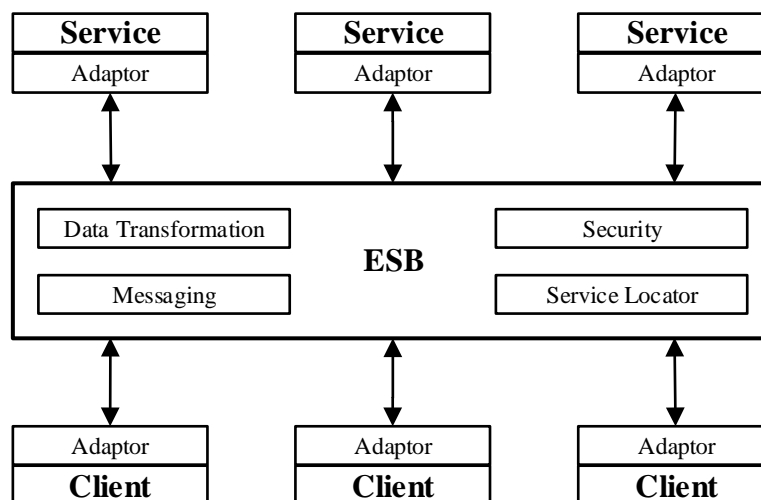


Figure 6.1: Architecture of an ESB

6.3 JBoss Fuse Middleware

6.4 ESB Application

6.4.1 Service Component Implementation

6.4.2 Service Component Logging

6.4.3 Service Component Tracing

Chapter 7

Design ESB in Openshift

As discussed in Chapter 6 on page 26, an ESB is a distributed computing architecture, where distributed services act as a consumer or producer. These services provide a business value in form of an integration of an internal or external service for an enterprise. There are multiple providers of ESB middleware like JBoss Fuse, as discussed in Section 6.3 on the preceding page, which provide tooling for implementing service components hosted on an ESB. It should be possible to migrate such a service component to a microservice, where features provided by the ESB middleware will have to be replaced by other implementations.

In this chapter an ESB application will be designed, where a service and a database will be integrated into each other by a another service. An Openshift Project will represent the service bus, which hosts the integration service, provides configuration and manages secrets for it. The concrete implementation of the services is considered to be not important, because they services will be ordinary Java web applications, which are known to be able to consume frameworks used in the Java Enterprise field. More important is the concept of how to represent a service component of an ESB as a microservice and how to host such microservices in Openshift which acts as the actual service bus.

7.1 Service Architecture

The Figure 7.1 on the next page illustrates the concept of the service architecture. This service architecture acts as an example of an application integration on an ESB. In an real world example such an service architecture would only represent a fraction of the actual services hosted on the ESB, which leads to the question how such an ESB can be monitored, especially when the integration services are hosted as separate microservices? The prototype will address the need for monitoring of the microservices by implementing tracing and logging features, which are discussed in Section 7.2 on the following page.



Figure 7.1: Service Architecture

7.1.1 Client

The client will be a Java REST-Client which consumes data from the backing service application. The REST-Client could be a single page web application or an other back-end application. It will have no direct access to the database or the integration service, which integrates the database and the Service application, and is therefore completely abstracted of the underlying data storage and its schema.

7.1.2 Service Application

The service application act as the back-end service for the client and produces data consumed by the client. The service application consumes from the integration service, which acts as the its back-end service. Same as the client, the service application is abstracted of the underlying data storage and its schema.

7.1.3 Integration Service

The integration service acts as the front-end service of the database and the back-end service for the consumers. The integration service ensures that the consumers are abstracted of the underlying data storage and its schema, as well as that the database is accessed in a proper manner, by providing a public API which defines the operations and models.

7.1.4 Database

The database acts as the data source for client which indirectly access this database via the integration service or a back-end service which itself is backed by the integration service. This level of abstraction of the consumers allows the database to evolve decoupled from the actual consumers. There is only the integration service which would have to be modified, when the underlying database or its schema changes.

7.2 Service Requirements

In this section the service requirements will be specified, which will ensure that the services are properly implemented and can effortlessly be managed and monitored.

Especially the management and monitoring becomes very important when moving from a conventional ESB application, like an application running on JBoss Fuse, to a microservice architecture, which is hosted on a PaaS platform like Openshift. Hosted on Openshift, the services run completely decoupled from each other with their own life cycle, which makes the management and monitoring harder compared to manage and monitor a monolithic application.

7.2.1 Technology

The services will be implemented in Java with the Java Enterprise Platform and the MicroProfile specifications. The MicroProfile specifications are an effort of the Eclipse Foundation to make Java applications ready for the cloud, where especially monitoring is a very important aspect to consider when it comes to distributed services. The services will be hosted as standalone applications like Spring Boot applications [Ecl17; Ecl18b].

7.2.2 State Handling

The service will be stateless, so that the services can be scaled and any request be handled by any instance of the service. Additionally Blue-Green-Releases and Canary-Releases are easily possible with stateless services [Fow10; Sat14]. Multiple instances of stateful service hosted on a PaaS platform are not flexible as stateless services, because sessions must stick to a particular service instance, and persistence volumes have to be shared between the service instances.

7.2.3 Monitoring

Monitoring is a essential aspect in distributed service architecture. Operators and developers need logging and tracing information of the services to comprehend errors in the services, where the cause of a problem could be located at another service as the service where a problem is reported.

Distributed Tracing

Distributed Tracing allows to comprehend service or method call chains. The MicroProfile specifications provides the OpenTracing specification, which provides an API for tracing an application on a method level or across service boundaries. The services must be able to collect tracing information about the REST and the related service method calls, and send this data to a central tracing service. The tracing implementation must be implemented separately from the service logic [Clo18b].

Distributed Logging

Distributed Logging allows to comprehend logs across service boundaries within a service call chain, where the logs of a service call chain have to be marked with a transaction id. The services must be able to provide all of their logging to a central service, whereby the logs are marked with a transaction id, which is equivalent to the transaction id of the service tracing. Optionally the services are allowed to add

additional markers, which can help developers and operators to analyze problems or to group service logs.

7.2.4 Configuration

The MicroProfile specifications provide the MicroProfile-Config specification, which provides an API to inject configuration parameters into classes, which can be loaded from different configuration sources. The services must provide the possibility to be configurable for different stages such as DEV (development), TEST (testing) and PROD (productive environment), where the services must not directly access the configuration sources, or hard code configuration in the source code unless its a default parameter [Ecl18a].

7.2.5 Fault Tolerance

The MicroProfile specifications provide the specification MicroProfile-Fault-Tolerance, which provides an API to define fault tolerance behavior such as retries, timeouts and error fall-backs. The fault tolerance of a service means that, if a depending service is not accessible at the time, a service must not fail immediately after the first try, but the service should retry to call the depending service for several times, and fail when all retries have failed. Such a behavior ensures that short timed communication errors, redeployments or overloads do not immediately cause a service to fail. The services must provide proper fault tolerance configuration and fall-back behavior to be able to recover from such errors in a proper manner. [Ecl18c].

7.2.6 API Management

The API management of a public API such as REST-API REST-Models ensure that the clients, using a public API, are not broken by changes made on that API. There are several opinions on how API management can be done. Swagger has become very popular for documenting REST-API, where the documentation can be used to generate clients, provide documentation for developers and to test the public API. The services must be capable of migrating their public API in a way that the clients are not broken by the change. The replaced API version must be supported as well as the new one, so that the client is not forced to modify its source code [Sma18].

ADD resource which christoph found during liwestfsw research

7.3 Openshift Architecture

7.3.1 Project

7.3.2 Templates

7.3.3 Scripts

7.4 Openshift Requirements

7.4.1 Monitoring

7.4.2 Configuration

7.4.3 Secrets

Chapter 8

Implementation ESB in Openshift

8.1 Technology

8.1.1 JBoss Fuse Integration Services

8.1.2 Wildfly Swarm

8.1.3 Fabric8

8.2 Monitoring

8.2.1 Service Implementation

8.2.2 Openshift Implementation

8.2.3 Jaeger

8.2.4 Kibana

8.3 Configuration

8.3.1 Service Implementation

8.3.2 Openshift Implementation

8.4 Fault Tolerance

8.4.1 Service Implementation

8.5 API Management

8.5.1 Service Implementation

8.6 Openshift

8.6.1 Configuration

8.6.2 Secrets

8.7 Implementation Microservice Architecture

which can be achieved with the framework Wildfly-Swarm. Wildfly-Swarm is the Java Enterprise answer to Spring Boot, where a Java Application Server is bootstrapped

programmatically, which only contains the minimal set of dependencies referenced by the service.

8.7.1 Distributed Tracing

8.7.2 Distributed Logging

8.7.3 API Management

8.8 Implementation Openshift Architecture

8.8.1 Security

Chapter 9

Evaluation ESB in Openshift

9.1 Software Development

9.2 Release Management

9.3 Security Handling

9.4 API Versioning

9.5 System Resources

References

Literature

- [Cor14] Intel Corporation. *Linux* Containers Streamline Virtualization and Complement Hypervisor-Based Virtual Machines*. Intel Corporation, 2014 (cit. on pp. 13, 14).
- [DG15] T. Devi and R. Ganesan. *Platform-as-a-Service (PaaS): Model and Security Issues*. School of Computing Science and Engineering, VIT University, 2015 (cit. on p. 1).
- [HW08] Gregor Hohpe and Bobby Woolf. *Enterprise Integration Patterns*. 11th ed. Boston, MA: Pearson Education, Inc., 2008 (cit. on p. 1).
- [Kie16] Kief Morris. *Infrastructure as Code: Managing Servers in the Cloud*. 1005 Gravenstein Highway North, Sebastopol, CA 95472: O'Reilly Media, Inc., 2016 (cit. on pp. 3–6).
- [Liu+15a] Lawrence Liu et al. *Integration Platform as a Service: The next generation of ESB, Part 1*. IBM, 2015 (cit. on pp. 1, 20).
- [Liu+15b] Lawrence Liu et al. *Integration Platform as a Service: The next generation of ESB, Part 2*. IBM, 2015 (cit. on p. 20).
- [Sch14] Mathijs Jeroen Scheepers. *Virtualization and Containerization of Application Infrastructure: A Comparison*. University of Twente, 2014 (cit. on p. 8).

Online sources

- [Ama18a] Amazon Web Services (AWS). *Amazon EKS*. 2018. URL: <https://aws.amazon.com/eks/> (visited on 03/31/2018) (cit. on p. 15).
- [Ama18b] Amazon Web Services (AWS). *AWS Elastic Beans Talk*. 2018. URL: <https://aws.amazon.com/elasticbeanstalk/> (visited on 04/12/2018) (cit. on p. 20).
- [Clo18a] Cloud Native Computing Foundation. *Sustaining and Integrating Open Source Technologies*. 2018. URL: <https://www.cncf.io/> (visited on 03/31/2018) (cit. on p. 15).

- [Clo18b] Cloud Native Computing Foundation. *Vendor-neutral APIs and instrumentation for distributed tracing*. 2018. URL: <http://opentracing.io/> (visited on 06/05/2018) (cit. on p. 30).
- [Cor18] CoreOS. *etcd*. 2018. URL: <https://coreos.com/etcd/docs/latest/> (visited on 04/01/2018) (cit. on p. 18).
- [Cos12] Costa, Glauber. *Resource Isolation: The Failure of Operating Systems and How We Can Fix It*. 2012. URL: <https://linuxconeuropa2012.sched.com/event/bf1a2818e908e3a534164b52d5b85bf1> (visited on 03/30/2018) (cit. on p. 12).
- [Doc18a] Docker Inc. *Docker Machine Overview*. 2018. URL: <https://docs.docker.com/machine/overview/> (visited on 03/17/2018) (cit. on p. 11).
- [Doc18b] Docker Inc. *Dockerfile Reference*. 2018. URL: <https://docs.docker.com/engine/reference/builder/> (visited on 03/09/2018) (cit. on p. 9).
- [Doc18c] Docker Inc. *Dockerfile Registry*. 2018. URL: <https://hub.docker.com/> (visited on 03/09/2018) (cit. on p. 9).
- [Doc18d] Docker Inc. *Overview of Docker Compose*. 2018. URL: <https://docs.docker.com/compose/overview/> (visited on 02/23/2018) (cit. on p. 6).
- [Doc18e] Docker Inc. *Swarm mode overview*. 2018. URL: <https://docs.docker.com/engine/swarm/> (visited on 03/30/2018) (cit. on p. 14).
- [Doc18f] Docker Inc. *What is Docker*. 2018. URL: <https://www.docker.com/what-docker> (visited on 02/23/2018) (cit. on p. 8).
- [Ecl17] Eclipse Foundation. *The Eclipse Enterprise for Java Project Top Level Project Charter*. 2017. URL: <https://projects.eclipse.org/projects/ee4j/charter> (visited on 06/05/2018) (cit. on p. 30).
- [Ecl18a] Eclipse Foundation. *Configuration for MicroProfile*. 2018. URL: <https://github.com/eclipse/microprofile-config> (visited on 06/05/2018) (cit. on p. 31).
- [Ecl18b] Eclipse Foundation. *Eclipse MicroProfile*. 2018. URL: <https://microprofile.io/> (visited on 06/05/2018) (cit. on p. 30).
- [Ecl18c] Eclipse Foundation. *Fault Tolerance*. 2018. URL: <https://github.com/eclipse/microprofile-fault-tolerance> (visited on 06/05/2018) (cit. on p. 31).
- [Fow10] Fowler, Martin. *BlueGreenDeployment*. 2010. URL: <https://martinfowler.com/bliki/BlueGreenDeployment.html> (visited on 06/07/2018) (cit. on p. 30).
- [Goo18a] Google Cloud. *Google App Engine*. 2018. URL: <https://cloud.google.com/appengine/> (visited on 04/12/2018) (cit. on p. 20).
- [Goo18b] Google Cloud. *Kubernetes Engine*. 2018. URL: <https://cloud.google.com/kubernetes-engine/> (visited on 03/31/2018) (cit. on p. 15).
- [Heo15] Heo, Tejun. *CGROUPS*. 2015. URL: <https://www.kernel.org/doc/Documentation/cgroup-v2.txt> (visited on 03/30/2018) (cit. on p. 13).

- [Kub18a] Kubernetes. *API Overview*. 2018. URL: <https://kubernetes.io/docs/reference/generated/kubernetes-api/v1.10/> (visited on 04/01/2018) (cit. on p. 17).
- [Kub18b] Kubernetes. *Kubernetes Components*. 2018. URL: <https://kubernetes.io/docs/concepts/overview/components/> (visited on 04/01/2018) (cit. on pp. 18, 19).
- [Kub18c] Kubernetes. *Pods*. 2018. URL: <https://kubernetes.io/docs/concepts/workloads/pods/pod/> (visited on 04/01/2018) (cit. on p. 17).
- [Kub18d] Kubernetes. *Production-Grade Container Orchestration*. 2018. URL: <https://kubernetes.io/> (visited on 03/30/2018) (cit. on pp. 14, 15).
- [Kub18e] Kubernetes. *Services*. 2018. URL: <https://kubernetes.io/docs/concepts/services-networking/service/> (visited on 04/01/2018) (cit. on p. 17).
- [Lin18a] Linux Containers. *Linux Containers*. 2018. URL: <https://linuxcontainers.org/> (visited on 03/29/2018) (cit. on p. 8).
- [Lin18b] Linux Foundation. *namespaces (7)*. 2018. URL: <http://man7.org/linux/man-pages/man7/namespaces.7.html> (visited on 03/30/2018) (cit. on p. 14).
- [Men18] Menage, Paul and Lameter, Christoph. *CGROUPS*. 2018. URL: <https://www.kernel.org/doc/Documentation/cgroup-v1/cgroups.txt> (visited on 03/30/2018) (cit. on p. 13).
- [Mic18a] Microsoft Azure. *Azure Cloud Services*. 2018. URL: <https://azure.microsoft.com/en-us/services/cloud-services/> (visited on 04/12/2018) (cit. on p. 20).
- [Mic18b] Microsoft Azure. *Azure Container Service (AKS)*. 2018. URL: <https://docs.microsoft.com/en-us/azure/aks/> (visited on 03/31/2018) (cit. on p. 15).
- [Ope14] Openshift. *OpenShift V3 Deep Dive Tutorial | The Next Generation of PaaS - Archived*. 2014. URL: <https://blog.openshift.com/openshift-v3-deep-dive-docker-kubernetes/> (visited on 04/05/2018) (cit. on p. 21).
- [Ope18a] Openshift. *App Development with OpenShift Online*. 2018. URL: <https://www.openshift.com/get-started/index.html> (visited on 04/12/2018) (cit. on p. 20).
- [Ope18b] Openshift. *OpenShift Application Platform*. 2018. URL: <https://github.com/openshift/origin> (visited on 04/12/2018) (cit. on p. 21).
- [Ope18c] Openshift. *Overview*. 2018. URL: https://docs.openshift.com/container-platform/3.5/architecture/core_concepts/ (visited on 04/07/2018) (cit. on p. 21).
- [Ope18d] Openshift. *Source-To-Image (S2I)*. 2018. URL: <https://github.com/openshift/source-to-image> (visited on 04/07/2018) (cit. on p. 24).
- [Red17] Red Hat Inc. *What is Wildfly?* 2017. URL: <http://wildfly.org/about/> (visited on 02/23/2018) (cit. on p. 6).

- [Red18a] Red Hat Inc. *Get Started with Red Hat JBoss Fuse 7 Tech Preview 3 Today!* 2018. URL: <https://developers.redhat.com/blog/2018/02/27/red-hat-jboss-fuse-7-tech-preview/> (visited on 04/12/2018) (cit. on p. 20).
- [Red18b] Red Hat Inc. *Red Hat JBoss Fuse*. 2018. URL: <https://developers.redhat.com/products/fuse/overview/> (visited on 02/22/2018) (cit. on p. 1).
- [Red18c] Red Hat Inc. *Red Hat OpenShift Container Platform*. 2018. URL: <http://www.openshift.com/container-platform/features.html> (visited on 02/22/2018) (cit. on p. 2).
- [Sat14] Sato, Danilo. *CanaryRelease*. 2014. URL: <https://martinfowler.com/bliki/CanaryRelease.html> (visited on 06/07/2018) (cit. on p. 30).
- [Sma18] SmartBear Software. *Swagger*. 2018. URL: <https://swagger.io/> (visited on 06/05/2018) (cit. on p. 31).